



Principais competências

Linux

Arquitetura de soluções

Modelagem de dados

Big Data & AI Specialist | Data Scientist & Engineer | Cloudera | Spark | Kafka | MLOps | LLMOps | GCP | Python | Master's in AI (UFG)

Curitiba, Paraná, Brasil

Resumo

Hi, I'm Vinícius — a Big Data & AI Specialist with 7+ years of experience in Data Engineering, Machine Learning, and Scalable Architectures. I'm currently pursuing a Master's in Computer Science with research focused on multi-agent architectures for LLMOps and RAG (Retrieval-Augmented Generation).

I work on the design and deployment of Big Data pipelines, secure Data Lakes, and MLOps/LLMOps architectures, leveraging tools like Cloudera, Hadoop, Spark, Kafka, and GCP. At Tecnisys and CEIA, I've been driving innovation in stream and batch processing, model interpretability, and production-ready AI applications.

My current focus is on:

"Architecting agent-based and multi-agent systems to optimize LLM orchestration and retrieval workflows."

"Designing secure data access layers using Data Mesh principles."

"Building scalable pipelines for time series prediction, NLP, and generative AI in Portuguese."

Highlights:

Delivered fraud detection and customer risk models at Copel using ML & NLP.

Implemented MLOps and LLMOps pipelines in production using GCP & open-source tools.

Leading research on AI interpretability, vector databases, and privacy-preserving RAG frameworks.

I'm passionate about turning complex data challenges into intelligent, scalable systems — and I'm always open to collaborative projects in AI, distributed systems, and cloud-native architectures.

Centro de Excelência em Inteligência Artificial (CEIA) & Deep Learning Brasil

Machine and Deep Learning Researcher

agosto de 2022 - Present (3 anos 11 meses)

I lead, design and implement enterprise-grade AI systems, bridging applied research and production environments.

- Lead R&D initiatives in Large Language Models, Agentic Architectures and Retrieval-Augmented Generation (RAG) systems for regulated and mission-critical environments.
- Architect and implement multi-agent LLM pipelines, ensuring scalability, governance and secure access control.
- Design and operationalize LLMops pipelines, including training workflows, evaluation frameworks, monitoring, retraining strategies and deployment automation.
- Implement Data Mesh-inspired architectures to enforce domain isolation, document-level access control and regulatory compliance.
- Deploy production AI systems on GCP, integrating custom infrastructure, open-source frameworks and scalable APIs.
- Lead cross-functional technical initiatives, aligning research, engineering and infrastructure teams to deliver AI solutions end-to-end.
- Build full-stack AI infrastructure: ingestion pipelines, vector databases, retraining loops, observability and performance optimization.
- Apply explainable AI (SHAP, XAI), time-series forecasting and statistical modeling in real enterprise use cases.

Tecnisys

Big Data Specialist

outubro de 2022 - Present (3 anos 9 meses)

I lead enterprise data architecture initiatives and implement large-scale distributed systems in high-demand environments.

- Architect, implement and administer enterprise Hadoop and Cloudera clusters.
- Lead Data Lake strategy and governance for batch and real-time workloads.
- Design and implement streaming and batch pipelines using Spark, Kafka, Hive and Python.

- Drive performance optimization, cluster tuning and distributed system resilience.
- Implement secure, production-ready pipelines with CI/CD, access control and infrastructure best practices.
- Coordinate cross-team technical decisions involving infrastructure, data engineering and application layers.

12 anos

Center of Excellence - Data Science, Big Data and AI

junho de 2020 - outubro de 2022 (2 anos 5 meses)

Technical leader in strategic AI initiatives within the energy sector.

- Led and implemented machine learning models for fraud detection, risk scoring and customer analytics.
- Architected and deployed scalable Big Data pipelines using Hive, Hadoop and PySpark.
- Integrated MLflow for experiment tracking, model governance and controlled deployment.
- Led PoCs from business conception to production deployment, engaging stakeholders across technical and executive levels.
- Spearheaded infrastructure specification for AI workloads, ensuring scalability and performance.
- Applied NLP, statistical modeling and ML techniques in operational and regulatory environments.

Java Web Systems Analysis, Design and Development

janeiro de 2019 - junho de 2020 (1 ano 6 meses)

- Developed data pipelines using Java and Kafka for ingestion and real-time processing.
- Created Web applications with JSF, Spring, and React, integrating them with Big Data systems.
- Gathered requirements, analyzed systems, and created documentation to support scalable, distributed solutions.
- Built Spark applications for efficient batch processing and ETL workflows.

Business Administrative Assistant

novembro de 2010 - janeiro de 2019 (8 anos 3 meses)

Formação acadêmica

Universidade Federal de Goiás

Mestrado em Ciências da Computação, Sistemas Inteligentes e Suas Aplicações · (janeiro de 2024 - agosto de 2026)

Universidade Federal do Paraná

Pós-graduação Lato Sensu - Especialização, Data Science e Big Data · (fevereiro de 2021 - junho de 2022)